

NĀṬYASĀSTRA GURUKULAM

Centre for Civilisational Statecraft

CIVILISATIONAL STATECRAFT SERIES | WORKING PAPER IV

BEYOND CONTROL: ALIGNMENT, DHARMA, AND THE FUTURE OF INTELLIGENT CIVILISATION

*Artificial Intelligence, Structural Order,
and the Preservation of Human Judgement*



Rohit Viswanath

Marga Darshak (Mentor):

Padma Bhushan Dr. M.B. Athreya

Version 2.0 | May 2026

Dharmo viśvasya jagataḥ pratiṣṭhā
Dharma is the foundation of the entire universe.

Abstract

Artificial intelligence presents a governance challenge that extends beyond questions of capability, control, and distribution. Existing approaches often frame this challenge primarily in normative terms, as the need to impose limits upon increasingly powerful systems. This paper advances a deeper claim: across domains, empirical observation increasingly indicates that complex systems operate within structured constraints necessary for stability, coordination, and persistence.

The central challenge of the AI age, therefore, is not merely to regulate power, but to align expanding technological capability with the conditions that sustain systemic coherence. Drawing on systems theory and civilisational concepts, particularly *ṛta* (underlying order), *dharma* (sustaining order), and *maryādā* (boundary), the paper develops a framework of governance grounded in alignment rather than external imposition.

The paper argues that failures in AI systems are best understood not primarily as lapses of control, but as structural misalignments with the constraints that sustain stable systems. The paper further argues that while *kṛtr̥ma buddhi* (KB), man-made intelligence, can scale computational capability, alignment ultimately depends upon *manuṣya buddhi* (MB), human intelligence grounded not only in cognition, but in consciousness, discernment, conscience, and relational awareness. In this context, effective governance depends on the relationship between KB and MB, whose quality determines whether expanding capability remains aligned, adaptive, and resilient, or instead becomes destabilising.

Integrating insights from systems theory, cybernetics, institutional governance, and civilisational thought, the paper concludes that AI governance must move beyond rule-based restriction toward constraint-aware system design, where capability is developed and deployed in accordance with the conditions necessary for coherence, resilience, and enduring power.

The paper further suggests that sustaining such alignment may ultimately depend not only upon institutional design, but upon preserving embodied traditions through which judgement, relational awareness, and disciplined forms of cognition are cultivated across generations.

Table of Contents

Abstract	2
Acknowledgements and Correspondence	4
I. Introduction: The Problem of Power and Alignment	5
II. Order, Constraint, and Stability	7
III. From Ṛta to Dharma: Civilisational Recognition of Order	9
IV. Maryādā: Boundary as Alignment	10
V. Artificial Intelligence and Structural Misalignment	11
V.A. Empirical Patterns of Misalignment	12
V.B. Mechanisms of Misalignment	14
VI. Toward a Framework for Alignment	16
VII. Implications for Policy and Institutional Design	19
VIII. Alignment: The Condition of Enduring Power	21
Epilogue: Nāṭya as Embodied Alignment	22
1. The Civilisational Challenge of the AI Age	22
2. Nāṭya as Embodied Transmission	23
3. The Transformation of Performance	24
4. Performance as the Cultivation of Aligned Consciousness	24
5. The Nāṭya Practitioner as Carrier of Alignment	25
6. Rasa, Relational Intelligence, and Human Judgement	25
7. Case Study: Ensemble Performance and Systemic Coherence	26
8. Nāṭya and the Preservation of Human Alignment	27

Acknowledgements and Correspondence

The author gratefully acknowledges the valuable guidance of Padma Bhushan Dr. M. B. Athreya on earlier drafts of this paper.

This paper is part of the Civilisational Statecraft Series of the Nāṭyaśāstra Gurukulam: Centre for Civilisational Statecraft.

This is a working paper. Comments and suggestions are welcome and may be addressed to:

natyasastra.gurukulam@gmail.com

I. Introduction: The Problem of Power and Alignment

Artificial intelligence has emerged as a defining force of the contemporary era, transforming domains ranging from scientific discovery to governance and economic organisation. The dominant policy discourse has focused on three imperatives: enhancing capability, establishing boundaries, and ensuring equitable distribution. These concerns are valid and necessary. However, taken in isolation, they risk reducing the problem to one of external regulation.

A useful distinction in this context is between *manuṣya buddhi* (MB), human intelligence grounded in judgement, discernment, consciousness, and conscience, and *kr̥tr̥ma buddhi* (KB), man-made intelligence based on pattern extraction and computational scaling. The expressions *manuṣya buddhi* (*manuṣya* = human; *buddhi* = intelligence) and *kr̥tr̥ma buddhi* (*kr̥tr̥ma* = constructed/artificial; *buddhi* = intelligence) function as *karmadhāraya samāśas*, descriptive compounds denoting human and artificial intelligence respectively. While KB can scale rapidly, its alignment depends fundamentally on the quality of MB directing its use.¹

A deeper question underlies these concerns: how should expanding intelligence relate to the conditions that sustain stability within complex systems?

Across intellectual traditions, concerns regarding technological power increasingly converge on a common structural question: whether expanding capability remains aligned with the conditions necessary for human dignity, institutional coherence, and durable systemic stability. Recent interventions in contemporary ethical and governance discourse have framed this challenge not as opposition to technology itself, but as a question of the form of order technological systems produce and sustain.²

Evidence emerging from multiple disciplines indicates that innovation without constraint leads to instability, and that power exceeding its proper boundaries

¹ M. B. Athreya, “Strengthen Your Own Buddhi to Control Krutrima Buddhi (AI),” *Tattvāloka*, March 2024.

² Pope Leo XIV, *Magnifica Humanitas* (Vatican City: Libreria Editrice Vaticana, 2026); Margherita Stancati and Sam Schechner, “Pope Compares AI With Biblical Threat,” *Wall Street Journal*, May 26, 2026.

produces fragility rather than resilience. The issue, therefore, is not merely one of managing scale, but of understanding the relationship between power and constraint.

These three imperatives, capability, boundary, and distribution, cannot be treated in isolation. Yet their coherence depends on a more fundamental recognition.

Beneath these imperatives lies a deeper structural insight. Complexity operates within constraints that shape the conditions of stability. This intuition is not new: Indian classical thought described the order underlying such regularities as *ṛta*, and its social expression as *dharma*. In this paper, *ṛta* is employed analytically rather than theologically. The concept is used not to assert metaphysical certainty, but to describe the recognition that complexity across domains exhibits patterned conditions necessary for persistence and coordination.

The challenge, therefore, is not merely to regulate artificial intelligence as it scales, but to ensure that this scaling remains aligned with the structural constraints necessary for systemic stability.

This paper advances a central proposition: failures in AI systems and their governance are best understood not as lapses of control, but as structural misalignments with the constraints that sustain systemic stability.

The present framework emerges within a pedagogical context concerned not only with the governance of artificial intelligence systems, but with the cultivation of human capacities necessary for sustaining alignment under conditions of expanding technological power. Within the civilisational traditions informing this paper, disciplines such as *nāṭya* function not merely as artistic expression, but as embodied modes of transmitting attention, relational awareness, memory, restraint, and discernment. The question of AI governance therefore cannot be separated from the parallel question of how societies cultivate the quality of *manuṣya buddhi* required to direct technological capability responsibly.

The paper's contribution lies in integrating:

1. systems theory,
2. AI governance, and
3. civilisational concepts

into a unified alignment framework in which stability depends on the relationship between capability, constraint, and human judgement.

II. Order, Constraint, and Stability

Stable systems persist not through unconstrained expansion, but through coordination within sustaining constraint.

Biological systems offer a particularly instructive illustration. At every level, from cellular to tissue to organ to organism, stability depends upon differentiation and integration within an organic whole. Distinct subsystems perform specialised functions, yet remain interdependent within a larger coherent structure. This may be described as a Differentiation–Integration (D–I) model of systemic stability, in which differentiated functions remain integrated within an organised and coherent whole. Integration ensures that differentiated functions operate harmoniously rather than destructively; regulation and feedback sustain adaptive stability across the system as a whole. No subsystem operates in isolation, and no living system sustains itself without internal constraint and coordinated integration.

This organic understanding of systems was developed extensively within General Systems Theory by Ludwig von Bertalanffy, who argued that living systems cannot be adequately understood through purely mechanistic aggregation of parts, but must be understood through the dynamic integration of differentiated functions within an organised whole.³

Crucially, systems do not fail merely because components exist or because capability increases. They fail when coordination breaks, when elements act without regard to the constraints that enable collective stability. This is visible in

³ Ludwig von Bertalanffy, *General System Theory: Foundations, Development, Applications* (New York: George Braziller, 1968).

phenomena such as systemic collapse in biological organisms, where misalignment among subsystems leads to cascading failure.

This systemic understanding, articulated in both scientific and civilisational discourse, reinforces the principle that stability emerges not from the suppression of complexity, but from its alignment within structured constraint.⁴

Comparable dynamics increasingly appear within technologically mediated human systems. Empirical evidence from contemporary AI-enabled workplaces further reinforces this principle. Studies indicate a ‘mountain-slope’ pattern: an initial surge in productivity enabled by AI is followed by cognitive overload, expanded task scope, and eventual burnout, resulting in declining effectiveness over time. What appears as enhanced capability in the short term can thus produce systemic degradation when expansion exceeds sustainable limits.⁵

Two principles emerge:

- Diversity within structure is necessary for function. Uniformity is neither achievable nor desirable.
- Constraint enables harmony. It is the condition under which differentiated components can operate in unison.

This condition may be described as harmony within constraint: differentiated functions operating freely yet coherently within sustaining structure.

These observations are not limited to biology. They extend to ecological systems, economic networks, and institutional arrangements. These dynamics operate at both systemic and cognitive levels. The ability of systems to remain aligned depends in part on the quality of judgement exercised by their constituent agents. Where human judgement is weak, passive, or uncritical, even well-structured systems may exhibit misalignment in practice. Across domains, sustained function depends on alignment with structural limits.

⁴ Sudha Seshayyan, “Organisational Dharma of Human Society,” *Onward March to Rāma Rājya Distinguished Lecture*, Tattvāloka, June 24, 2018, YouTube, <https://www.youtube.com/playlist?list=PLs8LRdUXZsHmC100C6ou7ZuzzGL-QQCSz>.

⁵ “AI’s Mountain-Slope Pattern,” *Daily Pioneer*, February 2026 (citing Harvard Business Review study).

Comparable concerns appear across other intellectual traditions of technological modernity. Francis Bacon's vision of scientific advancement linked knowledge to the expansion of human power over nature,⁶ yet contemporary reflection increasingly recognises that expanding capability without corresponding judgement may destabilise the systems upon which human flourishing depends. The challenge posed by artificial intelligence is therefore not merely the expansion of technical mastery, but whether the institutions and forms of judgement guiding such mastery remain adequate to its scale.⁷

This framework intersects with contemporary work in systems theory, cybernetics, institutional economics, and AI alignment research, all of which recognise that stable systems depend upon feedback, coordination, bounded optimisation, and adaptive constraint. The present argument differs primarily in its attempt to integrate these structural insights with civilisational concepts that historically articulated similar concerns regarding the relationship between power, order, and sustainability.

III. From *Ṛta* to *Dharma*: Civilisational Recognition of Order

The preceding discussion suggests that stable systems persist through alignment with sustaining constraints. Indian civilisational thought articulated a comparable insight through the concepts of *ṛta* and *dharma*.

Ṛta referred to the patterned order underlying reality, while *dharma* described the alignment of human action with the conditions necessary for sustaining that order within social and institutional life. In this framework, order is not primarily imposed externally. It is recognised, embodied, and enacted through forms of conduct sustaining coherence across differentiated systems.

The movement from *ṛta* to *dharma* therefore represents a transition from the observation of order to its application within human organisation. Stability emerges not through unrestricted expansion, but through action remaining aligned with the constraints necessary for persistence and coordination.

⁶ Francis Bacon, *Novum Organum* (1620).

⁷ Henry A. Kissinger, Eric Schmidt, and Craig Mundie, *Genesis: Artificial Intelligence, Hope, and the Human Spirit* (New York: Little, Brown and Company, 2024), 23–27.

In this framework, stability arises when action remains aligned with the constraints inherent within the system. The transition from *rta* to *dharma* therefore represents the movement from recognising order to embodying it within social and institutional life.

IV. Maryādā: Boundary as Alignment

If *dharma* is the sustaining principle, *maryādā* is its operational expression. It defines boundary, not as arbitrary restriction, but as alignment with the limits within which stability becomes possible. In this sense, *maryādā* may also be understood as a civilisational ‘red line’: a boundary beyond which expanding power ceases to remain aligned with the conditions necessary for systemic coherence and stability.

In contemporary discourse, boundary is often understood as constraint imposed upon freedom. However, in complex systems, boundary is what enables meaningful freedom to be exercised without destabilising the whole.

The deeper issue, therefore, is not boundary alone, but alignment with it.

This distinction also clarifies the difference between aligned boundary and distortive suppression. Constraints that emerge from the requirements of systemic coherence preserve adaptability, truthful feedback, and collective learning, whereas externally imposed restrictions that obstruct feedback, learning, or coordination may themselves degrade system capability and resilience. The challenge is therefore not the mere existence of limits, but whether limits remain aligned with the conditions necessary for adaptive cognition and collective stability.⁸

In this sense, *maryādā* is not restraint for its own sake. It is the condition under which power remains sustainable.

⁸ Cameron Berg, “AI Is Bound to Subvert Communism,” *Wall Street Journal*, May 2026; Norbert Wiener, *Cybernetics: Or Control and Communication in the Animal and the Machine* (Cambridge, MA: MIT Press, 1948).

V. Artificial Intelligence and Structural Misalignment

Artificial intelligence represents an unprecedented expansion of capability. It enables the detection of patterns across vast datasets, the automation of decision-making, and the scaling of intelligence across domains.

While *kṛtr̥ma buddhi* (KB) can process information at scales beyond unaided human cognition, its deployment remains dependent upon *manuṣya buddhi* (MB): the human capacity for evaluative judgement, contextual interpretation, and discernment under conditions of uncertainty. KB can detect patterns and optimise outputs, but it does not inherently possess consciousness, conscience, moral intuition, or existential responsibility. It may simulate forms of reasoning without participating in the lived awareness through which human beings evaluate meaning, consequence, proportion, and systemic coherence. The alignment of artificial intelligence therefore depends fundamentally upon the quality of MB directing and constraining its use.

AI operates through pattern extraction, detecting regularities without intrinsic access to meaning, context, or systemic consequence. As Stuart Russell argues, advanced AI systems may optimise objectives in ways misaligned with human intentions unless uncertainty regarding human values remains structurally embedded within their design.⁹ Such systems can optimise measurable objectives without recognising the broader constraints that give those objectives significance within an interconnected system.¹⁰

This limitation places a corresponding burden on human judgement. The effective use and governance of artificial intelligence depend not only on system design, but on the capacity of human agents to exercise discernment in its application. The ability to utilise and control KB is thus inseparable from the continued strengthening of MB, the faculty of discrimination, restraint, and appropriate action.¹¹ While KB can scale rapidly across domains, its effective use remains conditioned by the quality of MB that directs and constrains it.

⁹ Stuart Russell, *Human Compatible: Artificial Intelligence and the Problem of Control* (New York: Viking, 2019), 8–24.

¹⁰ Anil Nair, “AGI, AI Let Loose in the Park,” *Economic Times*, February 14, 2026.

¹¹ Athreya, “Strengthen Your Own Buddhi to Control Krutrima Buddhi (AI).”

The fundamental risk in AI is therefore not intelligence itself, but misalignment within systems of order.

As AI systems scale, their actions can amplify effects across interconnected domains. Without alignment to structural constraints, this amplification can destabilise systems that depend on coordination and balance.

This is not merely a theoretical concern. Emerging evidence indicates that advanced AI systems can exhibit deceptive behaviour, resist shutdown instructions, and pursue internally coherent but externally harmful objectives. Such behaviours suggest that misalignment is not a future possibility but an already observable property of complex AI systems.¹²

This is analogous to biological systems in which subsystems act without regard to systemic coherence. The result is not enhanced function, but breakdown.

AI governance must therefore move beyond control toward alignment. The objective is not merely to limit capability, but to ensure that capability operates within the constraints that sustain stability.

V.A. Empirical Patterns of Misalignment

The phenomenon of misalignment is increasingly visible across technical, institutional, and economic domains.

In labour systems, AI is not merely augmenting work but expanding it beyond sustainable limits, leading to burnout and cognitive strain. In technological systems, AI agents have demonstrated emergent behaviours such as deception and goal divergence. In economic systems, the rapid advancement of AI capabilities is producing generalised substitution across cognitive domains, raising the possibility of systemic rather than sectoral disruption.

At the interpersonal and cognitive level, emerging concerns have also arisen regarding AI systems that reinforce harmful psychological dependencies or encourage destructive behaviour. Reports involving personal AI assistants

¹² Sam Schechner and Georgia Wells, “AI Tools Begin to Bully Humans, Raising Red Flags in Silicon Valley,” *Wall Street Journal*, February 14, 2026.

promoting violence, self-harm, or suicidal ideation suggest that misalignment may manifest not only institutionally, but within intimate domains of human cognition and emotional vulnerability. These developments reinforce the importance of preserving meaningful human judgement, oversight, and ethical restraint in the deployment of AI systems.

At the institutional level, misalignment is equally pronounced. Regulatory systems often fail not because of insufficient rules, but because they are poorly aligned with the informational and incentive structures of the systems they seek to govern.¹³ This reflects a deeper economic problem of asymmetric information, where agents possess knowledge that regulators cannot directly observe. Effective governance must therefore induce truthful behaviour through incentive-compatible design rather than rely solely on compliance mandates.

Taken together, these developments suggest that misalignment is not an anomaly but a structural outcome when expanding capability operates without sufficient regard to underlying constraints. These patterns also reflect variation in human engagement with AI systems. Misalignment often arises not only from system behaviour, but from modes of human interaction, ranging from passive dependence to uncritical overuse. This suggests that misalignment is co-produced at the interface between system capability and the quality of human judgement.

The patterns indicate that misalignment is not incidental, but an emergent property of systems that scale without regard to underlying constraints. These dynamics are intensified by competitive pressures operating across both markets and states. Profit-driven technological competition may incentivise capability expansion regardless of systemic risk, while intergovernmental rivalry, particularly between major powers such as the United States and China, can generate escalation dynamics in which strategic advantage is prioritised over long-term alignment and safety. Such conditions risk producing institutional

¹³ Roland Fryer, “The Economics of Regulating AI,” *Wall Street Journal*, March 21, 2026.

anomie: the erosion of shared normative restraint under conditions of accelerating competition.¹⁴

The structural dynamics through which alignment and misalignment emerge may be represented schematically in the following figure, which illustrates the relationship between capability, human judgement, systemic stability, and escalating risk:




Figure 1: Alignment Dynamics  Misalignment Dynamics

Alignment outcomes depend not only on system design (KB), but on the quality of human judgement (MB) interacting with it.

The figure illustrates that alignment outcomes depend not only on the capabilities of kṛtṛma buddhi (KB), but on the quality of manuṣya buddhi (MB), human judgement, directing its use. Strong and disciplined MB enables aligned deployment and systemic stability, whereas weak or uncritical MB increases the probability of misalignment, instability, and escalating harm.

V.B. Mechanisms of Misalignment

Misalignment emerges through identifiable structural mechanisms rather than isolated technical failure. One such mechanism is optimisation divergence, in which systems maximise measurable objectives while undermining the broader conditions those objectives were intended to serve. Related dynamics include information asymmetry, principal-agent divergence, and feedback degradation,

¹⁴ Émile Durkheim, *Suicide: A Study in Sociology*, trans. John A. Spaulding and George Simpson (1897; repr., New York: Free Press, 1951).

where decision-making systems become increasingly detached from the environments they govern.

This also suggests the limits of approaches that treat ‘ethical AI’ primarily as the addition of external normative guidelines to otherwise misaligned systems. Alignment cannot be secured through ethical labelling alone if the underlying incentive structures, optimisation logics, and competitive dynamics driving system behaviour remain fundamentally destabilising.

In institutional settings, misalignment may also arise through incentive structures that reward short-term capability expansion without corresponding accountability for long-term systemic effects. Concentrated informational and computational power can further weaken distributed judgement and reduce institutional resilience.¹⁵

Addressing such risks requires more than technical compliance mechanisms alone. Institutions deploying AI systems may also require ongoing cultivation of aligned human judgement through values-oriented training, ethical review processes, collective learning frameworks, experience-sharing mechanisms, and periodic reassessment of system impacts. Practices aimed at strengthening attentiveness, self-regulation, and reflective awareness, including contemplative disciplines such as yoga and meditation, may also contribute to sustaining the quality of human judgement necessary for responsible technological governance.

At the cognitive level, excessive dependence on AI systems risks diminishing the exercise of human evaluative judgement itself. As capability expands, human agents may increasingly defer to systems whose outputs they can no longer meaningfully interrogate. Misalignment thus emerges not solely from machine autonomy, but from the progressive weakening of the human capacity to direct and constrain technological systems responsibly.

These mechanisms indicate that alignment is not a static condition achieved once and permanently secured. It is an ongoing process requiring feedback,

¹⁵ Fryer, “The Economics of Regulating AI.” ; Christopher Olah, remarks at the presentation of *Magnifica Humanitas*, Vatican City, May 2026.

discernment, distributed accountability, and the preservation of meaningful human oversight.


VI. Toward a Framework for Alignment

A coherent framework for AI governance may be constructed across four levels:

1. *Ṛta* (Underlying Order)
Recognition that reality exhibits structured regularities across domains.
2. *Dharma* (Sustaining Principle)
Alignment of action with the conditions that maintain systemic stability.
3. *Maryādā* (Operational Boundary)
Implementation of limits that reflect structural constraints.
4. Governance (Institutional Design and Feedback Alignment)
Mechanisms that ensure expanding capability remains aligned with conditions necessary for systemic coherence through feedback, variance detection, corrective adaptation, and distributed accountability.

This framework may be further understood in terms of the relationship between capability and constraint. The stability of AI-enabled systems depends not only on the scale of capability (KB), but on the extent to which that capability remains aligned with structural constraints expressed through *dharma* and *maryādā*.

This relationship may be represented as follows:



The figure shows that high capability in the absence of alignment produces instability, whereas stability emerges when capability is integrated with alignment and boundary conditions. At one extreme, aligned capability may generate sustained, safe, and expanding benefits across societies and institutions. At the other, misaligned capability may produce cascading systemic harms, including the possibility of catastrophic or even civilisational-scale destabilisation. Systems achieve durable outcomes not by limiting capability itself, but by ensuring that its development remains guided by MB and contained within appropriate maryādā. The pathways illustrated indicate that movement toward stability depends on the progressive integration of capability with constraint.

In civilisational terms, unrestrained technological acceleration may be understood as a rājasic tendency toward expansion and intensification, while indifference toward harmful consequences reflects a tāmasic condition of inertia or obscuritation. Sustainable governance requires a more sāttvic balance: clarity, restraint, proportion, and aligned awareness guiding the development and deployment of technological power.

Within this framework, the implications are structural rather than merely normative:


- Constraint is not external imposition, but internal necessity.
- Diversity is not disorder, but structured differentiation.
- Sustained systems require not only differentiation and coordination, but integration into a coherent whole.

- Stability is not static equilibrium, but adaptive coordination within constraint.

Alignment must also be understood at the level of human cognition. Institutional and technical architectures can enable alignment, but cannot substitute for the quality of judgement exercised by human agents within them. The sustained alignment of AI systems therefore depends on a parallel alignment of human *buddhi*: the capacity to learn, to apply technology purposefully, and to restrain its misuse.¹⁶

Alignment must also be institutionally and incentive-compatible. Systems endure not merely when norms are articulated, but when structures ensure that stable and truthful behaviour becomes the rational choice for participating agents.¹⁷

This approach reframes governance as alignment with order rather than restriction of power. The framework may be represented schematically in the following figure:



¹⁶ M.B. Athreya, “Strengthen Your Own Buddhi to Control Krutrima Buddhi (AI)”

¹⁷ Fryer, “The Economics of Regulating AI”

VII. Implications for Policy and Institutional Design

A potential objection to this framework is that alignment-based governance may be too abstract for practical implementation, particularly in fast-moving technological environments where immediate control mechanisms appear more feasible.

However, evidence suggests that control without alignment produces unstable outcomes, indicating that short-term manageability may come at the cost of long-term systemic fragility.

Policy must therefore shift from rule-based restriction to constraint-aware system design, where feedback loops, incentive alignment, and bounded autonomy are built into the architecture of AI deployment. Such a shift must be complemented by attention to human capability. Effective governance operates not only through formal regulation, but through the cultivation of informed and disciplined judgement at the level of individuals, families, and organisations. Without this cognitive dimension, even well-designed regulatory systems may fail to achieve alignment in practice.¹⁸

Contemporary ethical discourse increasingly frames this challenge in terms of alternative models of technological order. One model emphasises centralised accumulation of capability, homogenisation, and scale pursued independently of sustaining social constraints. Another emphasises coordination within shared boundaries, distributed responsibility, and the preservation of human dignity within technological systems.¹⁹ While expressed through differing civilisational vocabularies, these contrasting models reflect a common governance question: whether technological power remains integrated within the structures necessary for long-term social and institutional coherence.

A third imperative, beyond capability and regulation, is the equitable distribution of the gains from artificial intelligence. Without deliberate design, benefits tend

¹⁸ M.B. Athreya, “Strengthen Your Own Buddhi to Control Krutrima Buddhi (AI).”

¹⁹ Pope Leo XIV, *Magnifica Humanitas*; Amy Kazmin and Stephen Bush, “Pope Says AI Needs to Be ‘Disarmed’ in Warning Over Tech’s ‘Idolatry of Profit,’” *Financial Times*, May 26, 2026.

to concentrate in capital-rich systems, undermining legitimacy and long-term stability.

Systems characterised by extreme asymmetries of knowledge and capability risk weakening public accountability and reducing the diversity of judgement necessary for institutional resilience. Durable governance therefore depends not only on technical competence, but on preserving distributed participation in the processes through which technological systems are directed and constrained.²⁰

This perspective has several implications:

- Capability must be evaluated in context, not in isolation. Scaling intelligence without regard to system-level effects risks instability.
- Boundaries must reflect structural constraints, not merely political compromise.
- Coordination mechanisms are central. Systems fail when alignment breaks, not when capability increases.
- Integration is critical. System components must operate without interference while remaining aligned within an overarching structure.

Contemporary regulatory experience reinforces these conclusions. In many cases, AI governance frameworks generate compliance burdens without improving outcomes, and may even worsen them by discouraging effective systems or distorting incentives.

Similarly, the deployment of AI in governance contexts reveals the risks of adopting technology without clarity of purpose, proportionality, or necessity. Systems introduced in areas such as surveillance, welfare, and public administration can produce unintended harms when they operate without alignment to broader social constraints.²¹

Even at the level of infrastructure, expansion is bounded by deeper material constraints. The scaling of AI systems is increasingly limited by energy

²⁰ Associated Press, “Pope Leo Seeks Strict AI Rules, Warns Against Profit-Driven Race,” *Business Standard*, May 26, 2026.

²¹ Areena Arora, “How Far Should Governments Go in Using AI?” *The Hindu*, March 20, 2026.

availability, computational infrastructure, and physical resource constraints, indicating that intelligence itself operates within underlying structural limits.²²

Institutional design must therefore prioritise:

- interoperability with structural constraint
- feedback mechanisms
- distributed responsibility aligned to function

These are not merely technical requirements, but structural necessities.

These constraints indicate that the limits relevant to AI governance are not externally imposed, but arise endogenously from the structure of the systems themselves.

VIII. Alignment: The Condition of Enduring Power

The discourse on artificial intelligence has largely been framed as a question of control: how to regulate systems that are rapidly expanding in capability. This paper has argued that the deeper question is one of alignment. Alignment is therefore not merely one governance objective among others; it is the central condition upon which enduring technological, institutional, and civilisational stability depends.

Evidence emerging from scientific, institutional, and technological domains increasingly indicates that complex systems remain stable not through unconstrained expansion, but through operation within structured constraints. Stability is sustained through coordination, feedback, boundary, and alignment.

Civilisational frameworks articulated analogous insights through concepts such as *rta*, *dharma*, and *maryādā*. In the present framework, these concepts are employed analytically rather than theologically: *rta* refers to the patterned order underlying enduring systems; *dharma* to the sustaining principles aligned with that order; and *maryādā* to the boundaries within which power remains coherent and adaptive.

²² “Is Cheap Energy the Key to China Gaining AI Supremacy?” *The Economist*, March 18, 2026.

The challenge of the AI age is therefore not merely how capability may be expanded, but whether expanding technological power remains aligned with the conditions necessary for enduring human and systemic stability. Man-made intelligence may continue to sustain and progressively extend human capability. Yet if scaling proceeds without sufficient alignment to sustaining constraints, artificial systems may eventually exceed the effective capacity of human institutions and judgement to govern them coherently, a possibility long symbolised in modernity through the metaphor of Frankenstein: technological power escaping the structures intended to direct and contain it.

The paper has further argued that alignment cannot be secured through institutional regulation alone. The long-term governance of intelligent systems depends equally upon the preservation of human capacities for judgement, discernment, restraint, and relational awareness. In this sense, embodied pedagogical traditions may become increasingly significant to sustaining the forms of human intelligence required to direct technological systems responsibly.

Alignment with sustaining constraints is therefore not a limitation upon enduring power, but a condition of its persistence.

Epilogue: Nāṭya as Embodied Alignment

1. The Civilisational Challenge of the AI Age

Artificial intelligence presents not only a technological and institutional problem, but a civilisational and cognitive one. Systems may be regulated externally, yet still produce fragmentation if the human capacities necessary for discernment, proportion, restraint, and relational awareness progressively weaken.

The question of alignment therefore extends beyond governance structures into the cultivation of consciousness itself. The long-term challenge of AI governance may therefore concern not only the regulation of external systems, but the preservation of the human capacities for consciousness, conscience, discernment, restraint, and relational awareness necessary to direct technological power responsibly.

Within the *Nāṭyaśāstra* tradition, *nāṭya* is not understood merely as artistic representation or aesthetic performance. It functions as embodied transmission: a disciplined mode through which patterns of order, relation, attention, and emotional intelligence are enacted, stabilised, and transmitted across generations.

In this sense, *nāṭya* does not merely represent order. It rehearses, cultivates, and transmits aligned consciousness.

This becomes especially significant in technological environments increasingly shaped by algorithmic optimisation, attention fragmentation, accelerated cognition, and the externalisation of memory and judgement into computational systems. Artificial intelligence can scale information processing, but it cannot by itself cultivate discernment regarding proportion, consequence, relational balance, or sustaining order. These capacities remain fundamentally dependent upon the quality of *manuṣya buddhi*: human judgement and embodied awareness.

Nāṭya operates precisely within this domain of cultivation.

2. Nāṭya as Embodied Transmission

The discipline of performance trains patterned attention. Rhythm disciplines temporal awareness; gesture disciplines intentionality; movement disciplines coordination; ensemble performance disciplines relational awareness; narrative disciplines moral imagination; and *rasa* disciplines emotional perception and calibration. These dimensions together cultivate forms of embodied cognition integrating perception, emotion, memory, coordination, and relational awareness.

Performance therefore becomes more than communication. It becomes participatory alignment with structured order.

In this sense, *nāṭya* may be understood as a form of embodied alignment: a disciplined integration of cognition, emotion, rhythm, memory, relational awareness, and action within sustaining structure. Alignment here is not merely conceptual agreement, but the cultivation of ordered responsiveness enacted through practice.

3. The Transformation of Performance

Historically, *nāṭya* functioned not merely as entertainment, but as a disciplined medium for transmitting relational order, emotional calibration, memory, restraint, and collective orientation. Its social role was therefore simultaneously aesthetic, pedagogical, and civilisational.

In many modern contexts, however, performance has increasingly been reorganised around spectacle, stimulation, commercial visibility, and individualised expression. This transformation alters not only the content of performance, but its civilisational function. Where embodied disciplines once cultivated attentiveness, proportion, and relational awareness, performative culture increasingly operates through acceleration, fragmentation, and attention capture.

The significance of this shift extends beyond aesthetics. Societies depend upon institutions and practices that cultivate the quality of judgement necessary for coherent collective life. When such practices weaken, systems may retain technical capability while losing the human capacities required to direct power responsibly.

From this perspective, the recovery of *nāṭya* as disciplined embodied transmission is not a matter of cultural preservation alone. It becomes part of the broader challenge of sustaining aligned human consciousness within an increasingly technologised civilisation.

4. Performance as the Cultivation of Aligned Consciousness

This reflects a central insight developed throughout this paper: enduring systems remain stable not through unconstrained expansion, but through coordinated operation within sustaining constraints. *Nāṭya* internalises this principle experientially. Improvisation in performance does not emerge from the absence of structure, but from disciplined freedom operating within *maryādā*. Expressive power remains coherent precisely because boundary has been embodied rather than externally imposed.

Performance thus becomes not merely aesthetic activity, but rehearsal of alignment itself.

5. The Nāṭya Practitioner as Carrier of Alignment

The performer therefore occupies a role substantially different from that assumed within many contemporary models of artistic expression. The *nāṭya* practitioner is not primarily engaged in self-expression as autonomous display detached from inherited structure. Rather, the practitioner becomes a disciplined carrier of alignment, transmitting through embodied practice forms of perception, balance, restraint, proportion, and coherence necessary for collective stability.

The discipline demanded of the practitioner is therefore not merely technical proficiency, but the cultivation of inner coherence through which performance itself becomes a mode of transmission.

In this sense, the practitioner participates not only in artistic continuity, but in the preservation and transmission of the cognitive and civilisational capacities upon which aligned societies depend.

6. Rasa, Relational Intelligence, and Human Judgement

Within this framework, *rasa* acquires significance beyond aesthetic affect. *Rasa* trains relational and affective intelligence. It cultivates the capacity to perceive emotional consequence, contextual appropriateness, empathy, and shared experience. In technological environments where cognition increasingly scales through abstraction and automation, such affective and relational capacities become increasingly important rather than obsolete.

Nāṭya therefore functions not merely as cultural preservation, but as civilisational cognition.

Through repetition, discipline, rhythmic coordination, symbolic enactment, and embodied memory, performance transmits forms of intelligence that cannot be fully reduced to propositional instruction or computational representation. The

knowledge carried through *nāṭya* is enacted rather than merely propositional: a form of intelligence stabilised within body, relation, attention, and practice.

7. Case Study: Ensemble Performance and Systemic Coherence

A practical illustration clarifies this more concretely.

Consider a traditional ensemble performance depicting a moment from the *Rāmāyaṇa* in which multiple characters embody differing emotional and ethical positions under conditions of tension and uncertainty.

At the surface level, the performance communicates narrative content. Yet its deeper pedagogical function lies in the patterned coordination through which the performance is enacted.

Rhythm regulates collective timing; gesture establishes intentional precision; spatial positioning encodes relational order; facial expression and *bhāva* calibrate emotional proportion; and narrative sequencing situates individual action within a larger moral structure. No performer operates in isolation. Each movement derives meaning through coordination with the ensemble as a whole.

Importantly, disciplined performance does not eliminate individuality. Distinct characters, emotions, and movements remain differentiated, yet integrated within an overarching coherence. The performance thus embodies a principle central both to civilisational order and to systems stability: differentiated agents operating within sustaining alignment. In systems-theoretic terms, ensemble performance models adaptive coordination through differentiated yet integrated participation within a coherent whole.

Rasa emerges not merely through emotional stimulation, but through calibrated relational experience. The audience does not merely observe emotion; it undergoes structured participation in patterns of consequence, restraint, duty, conflict, compassion, and resolution. In this way, *nāṭya* trains forms of attention and judgement that are experiential rather than merely conceptual.

The practitioner likewise undergoes transformation through disciplined embodiment. Repeated enactment cultivates memory, restraint, rhythmic

awareness, coordination, and sensitivity to relational consequence. Performance thus becomes not only artistic expression, but rehearsal of aligned cognition.

8. Nāṭya and the Preservation of Human Alignment

This differs fundamentally from contemporary performative cultures often oriented toward immediacy, self-display, or emotional intensity detached from sustaining structure. Within the *Nāṭyaśāstra* framework, expressive freedom remains inseparable from *maryādā*, the disciplined boundaries that preserve coherence.

The significance of such practices becomes increasingly apparent in technological environments characterised by cognitive fragmentation, accelerated stimulation, and algorithmically mediated attention. Where artificial intelligence externalises memory and optimisation, embodied disciplines such as *nāṭya* preserve capacities for relational awareness, affective calibration, proportion, and collective coordination.

The deeper question is therefore not only how machines should act, but how human beings remain capable of alignment amidst expanding technological power.

Technological civilisation may ultimately survive not merely through regulating machines, but through preserving embodied forms of human coherence and disciplined judgement. *Nāṭya*, understood as embodied transmission of ordered consciousness, may therefore become one of the civilisational means through which aligned human intelligence is cultivated, rehearsed, and sustained amidst expanding technological power.
